



Full Stack Web Application Performance Tuning

codecentric GmbH

About the presenter





Fabian Lange, 27

married, no children but 2 cats

Fabian & symfony

Fabian.Lange@symfony-project.com

Using symfony since late 2006

Core team member since early 2008

Release manager for symfony 1.2

Fabian & work

Working as Java & Web architecture and performance consultant at codecentric

www.codecentric.de

Previously worked for Nokia / Nokia Siemens Networks and Hewlett Packard

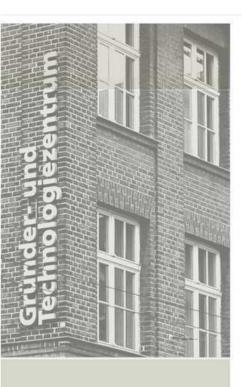
Fabian & blogging

Company blog at <u>blog.codecentric.de</u>

Private blog at blog.hma-info.de

About codecentric





Performance

Specialized in Java Performance optimization, codecentric has 100% success rate in finding and solving performance issues at a huge number of customers. We strive to continue this rate using state of the art tools and processes.

Architecture

Consulting software architects or managers, assisting with the design and creation of modern software architecture, weather it shall be a SOA architecture or classical design, our experts will get you there.

Open Technology

Cost efficient Open Source Tools and Solutions complete our portfolio. We believe that using Open Technologies improve software products quality and are more cost efficient than reinventing the wheel.

www.codecentric.de

See our full portfolio, meet the team and get our contact details.

Scope of this presentation

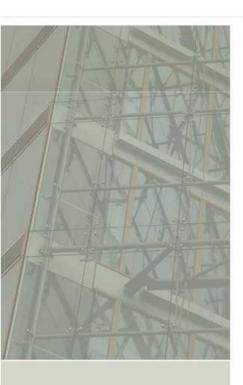




- Make web applications feel faster
 - Ideally a stopwatch can prove out effort
 Performance is perceived differently
 Tweak at the end of development cycle
- Not an expert guide on specific issues
 - But see the whole picture No micro-tweaking
- Targeted at small to medium deployments
 - Larger ones have dedicated people or high end hardware
 - Application developer tweaks single server
 - Keep scaling options open share nothing

What is the Whole Picture?





- Client Software Browser
- Web Framework symfony
- Server Software

PHP

Web Server

Database

Server Hardware

Processor

Memory

Storage

Client Software – Browsers





Rendering priority

Javascripts at bottom

Can be done by explicitly outputting them in the template
<?php echo include_javascripts();?>
To make this possible avoid inline Javascript or use Event.observe
Watch out with using JS before it is loaded
For example mouseover events
Prototype Event.observe as well!

CSS at top

Layout can be already calculated when building the DOM tree.

Image dimensions

Same as with CSS. Browser knows the image size before it loads the image Not only for performance, but flickering UI looks bad as well.

Client Software – Browsers





- Obtaining HTML is often the smallest part of the browser latency
- Limited connections (2-6) additionally slow down loading

Use a separate host for assets

You can even use another server for your media files

Even better try consider using a CDN -> less load

Cookies are not sent to other domains, which reduces request size

Get rid of unneeded additional requests

Many pages use a lot of images

```
CSS Sprites
```

```
#nav li a {background-image:url('/img/image_nav.gif')}
#nav li a.item1 {background-position:0px 0px}
#nav li a.item2 {background-position:0px -32px}
```

Inline Image Data (RFC2397)

```
<img src="data:image/gif;base64, ASCIIDATA" />
```

Client Software – Browsers





Get rid of unneeded additional requests (2)

Combine JS & CSS

Automation tools available

Sensible modularization allows separate update cycles and selective inclusion Protoculous: 1 file, but pages that just want Prototype have to parse Scriptaculous http://code.google.com/apis/ajaxlibs/

Minify JS and CSS

Easy and compatible using YUI Compressor

Dean Edwards Packer minifies better but produces runtime overhead

Use expires for assets

Note: You need to change file name for users to see the change

ExpiresActive On

ExpiresDefault "access plus 10 years"

Use redirects wisely

Redirects make sense to prevent double postings
Redirects are responses to the browser with a instruction to load a new page





Performance Tweaks

ETags

used to save bandwidth, but does not help with latency calculates a hash code of the returned content and only returns content if doesn't match browsers Etag doesn't work well with external compression

Core Compilation

symfony creates a single PHP files with all its core classes

Comments are removed from this file

Own classes can be added via app/config/core_compile.yml

- %SF_ROOT_DIR%/apps/frontend/lib/myUser.class.php
- %SF_ROOT_DIR%/lib/AutoLinker.class.php

Should be done only with classes





Propel

Use Propel 1.3 instead of 1.2 as it is faster and less memory consuming

Mainly benefits from object instance pooling

Monitor "slow queries" log

Use Explain \$slow_query

Add index where no index used

SELECT section.ID, section.TITLE, section.TREE_LEFT,
section.TREE_RIGHT, section.TREE_PARENT, section.TREE_SCOPE FROM
`section` WHERE section.TITLE='_news' LIMIT 1;

select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
SIMPLE	section	ALL	NULL	NULL	NULL	NULL	2121	Using where
	A a la la	•	possible_keys	key	key_len	ref		Extra
SCIECT TVNC	Tanie	TVDE	nossinie kevs	KAV	kev len	ret	rows	FYTTA
sciect_type	tubic	cypc.	possible_keys	KC y	Key_icii		10443	LATIC

Use Peer:doSelectJoinxxx when many related objects are loaded Reduce unneeded hydration

setUserID(\$comment->getUserId()) instead of setUser(\$comment->getUser())

Reduce unneeded queries

DB access is the most expensive part, it is always good idea to use few queries.

In extreme cases consider tweaking the database model by moving columns and merging tables, as a perfectly normalized database model is slow due to many joins.





Caching

Caching is very easy!

Create a cache.yml in the config directory of a module
List the actions/components/partials you want to cache
action:
enabled: on

Two important options:

with_layout contextual

There are more cacheable partials than expected

Caching nearly always improves response times

Don't build complex cleaning

Introduces risk of not cleaning everything

Clean every hour

Clean more than needed

New wildcard cleaning in symfony 1.1





symfony is modular

It is easy to bypass parts

It is easy to exchange parts

```
#app/config/factories.yml
all:
   storage:
    class: myMemorySessionStorage
```





Session Storage

PHP stores user sessions by default in files symfony does the same by using sfSessionStorage Files on the filesystem can be slow, can be a security risk and are hard to scale symfony ships alternative Session Storage classes

Note: reusing a propel connection for MySQL does not work

Server Software





Server Software often differs, especially in versions

PHP

PHP itself gains performance boosts with every release -> later is better http://sebastian-bergmann.de/archives/745-Benchmark-of-PHP-Branches-3.0-through-5.3-CVS.html

Web Server

Both Apache HTTPD and LightHTTPD work fine

http://www.markround.com/archives/30-LigHTTPd-and-Apache-Symfony-benchmarks.html
Using PHP via FastCGI brings huge performance boost!

But brings also troubles, e.g. for APC user cache

Database

A standard setup usually ships a MySQL

InnoDB should be preferred because of its tuning options and reliability

MyISAM is claimed to be "faster" but in fact that is rarely the case.

Of course other DBs are good as well ©

codecentric

Server Software – PHP





- As PHP becomes faster with each version, there should be no real reason to stay with old versions
- PHP itself allows not much performance tweaking
- A PHP opcode compiler should be mandatory, and actually will be included in PHP 6 by default

```
APC
    extension=apc.so
    apc.stat=0
        Disables APC checking for PHP file modifications
        Requires a server restart to detect updated files
xCache
    extension = xcache.so
```

Server Software – Apache 2.2





Multi Processing Module

Prefork

Will launch Apache processes that launch multiple PHP instances

Worker

Will launch Apache processes that launch one PHP instance with multiple threads

Worker works fine with most symfony applications, but some modules might not work

Logging

Frequently accessed resources will cause excessive logging.

To selectively turn off logging use this:

```
#apache2.conf
    CustomLog /var/log/apache2/access.log combined env=!dontlog
#Site/v-host config
    SetEnvIf Request_URI "^/chat" dontlog
```

Or perhaps turn off logging at all ©

Server Software – Apache 2.2





Compression

There are 3 methods of compression possible:

Let Apache compress
Let PHP compress

Let symfony compress

Apache compression is most proven and easy to configure

AddOutputFilterByType DEFLATE text/html text/plain text/xml text/css application/x-javascript works also on non PHP & non symfony content

Make sure to only compress with one option

Other Tweaks

HostnameLookups off

will prevent reverse lookup of hostename for IPs

ExtendedStatus off

will prevent apache collecting too much status information

Server Software – Apache 2.2





Keep alive

This HTTP 1.1 feature works roughly like this

When a browser connects to load a page a connection is opened

Instead of closing the connection after returning, lets say the HTML, with keep-alive the connection is kept open.

When one of two criteria is reached the connection is closed. This can be either a number of requests, or a timeout.

The upside of this feature is that some connection overhead is reduced

The downside is that resources are wasted in idle periods

Often the browser will load images directly after the HTML, unless the are cached

Many guides recommend to turn this feature off, but for small to medium loaded servers this can improve the responsiveness

A balanced config could look like

```
KeepAlive On
MaxKeepAliveRequests 1000
KeepAliveTimeout 3
symfony 1.0 - Need to patch symfony response class
sfWebResponse::sendHttpHeaders()
HTTP/1.0 -> HTTP/1.1
```

Server Software – MySQL 5





MySQL configuration done in config file

/etc/mysql/my.cnf

MySQL itself has multiple sample configurations, check them

Generic MySQL settings

 $key_buffer = 32M$

MyISAM but also for temp tables

thread_cache_size = 32

Creating threads is slow, try to get "threads created" low

table_cache = 1K

Opening tables is slow

sort_buffer_size = 16M

Used for sorting results, increase if "sort_merge_passes" increases

query_cache_limit = 1M

Caches how queries are executed, does not bring much but also does not cost much

query_cache_size = 32M

Caches results of queries, but too high settings involve high maintenance costs

Server Software – MySQL 5 & InnoDB





Disable Logging

$$\# \log = ...$$

Specify Slow Logging for Optimisation

```
log_slow_queries = /var/log/mysql/mysql-slow.log
long_query_time = 2
log-queries-not-using-indexes
```

InnoDB settings

```
innodb_log_file_size = 250M
    Data written goes into here first
innodb_buffer_pool_size = 1G
    Main InnoDB data cache. Waste memory here!
innodb_additional_mem_pool_size = 20M
    Utility buffer for InnoDB
innodb_thread_concurrency = 8
innodb_flush_log_at_trx_commit = 2
    2 will flush to OS cache, 1 will write to disc every commit
```

Server Hardware





Processors

More are better, but scales only until I/O wait Even the most basic server offerings have a dual core by now

Memory

Reduces I/O wait, as data can be in memory

An Apache process running symfony needs up to 64 MB

Also useful for database caching

Think of: DB Size x 2

Storage

Web servers mostly spend time on loading database, html or asset data Get fast discs. Noise doesn't matter in server rooms Use at least a two disc setup, if no raid

Main hardware issue is I/O wait!

Server Hardware – Two Disc Setup





- Even without expensive and perhaps even complicated raid configurations a second harddisc might significantly boost the server performance
- Separate discs by responsibility

Example:

One user does a heavy DB operation

In meantime 10 other users want to browse the page, requesting 10 images each

One disc can read/write for the DB operation continuously

The other disc can handle the many small read bursts for images

Simple config for InnoDB:

```
innodb_log_group_home_dir = /var/disc1/log/mysql/iblogs
innodb_data_home_dir = /var/disc2/mysql
innodb_data_file_path = ibdata:2000M:autoextend
```

References





Symfony

http://www.symfony-project.org/book/1 1/18-Performance http://trac.symfony-project.org/wiki/Optimization

Server:

http://typo3.org/development/articles/using-php-with-mod-fcgid/http://httpd.apache.org/docs/2.2/misc/perf-tuning.html
http://www.devside.net/articles/apache-performance-tuning
http://www.mysqlperformanceblog.com

CSS & JavaScript

http://www.alistapart.com/articles/sprites
http://developer.yahoo.com/performance/

http://developer.yahoo.com/yui/compressor/

http://developer.yahoo.com/yslow/





codecentric GmbH

Grünewalder Str. 29-31 42657 Solingen

phone +49-212-2494315 fax +49-212-2494109 email info@codecentric.de

codecentric